
Project guide

Python-only reproduction package for the MAS LCE fraudulent-classified-ads study

March 2026

What this package is

This package is a self-contained English compendium built from the archived David Graz MAS LCE project. It contains:

- an article-style reconstruction of the thesis;
- a dedicated integrity and reproducibility audit;
- sanitised canonical datasets that are sufficient to reproduce the published descriptive tables and the archived predictive metrics;
- Python scripts that regenerate tables, figures, manifests, checksums, and PDF reports;
- metadata files that document provenance, version drift, and canonicalisation choices.

1 Package structure

Path	Contents
reports/	Final PDFs and LaTeX sources for the article, audit, and guide
figures/	Reproduced figures used by the reports
data/canonical/	Sanitised canonical exploratory and case-study datasets
data/tables/	CSV tables used in the reports
data/json/	JSON summaries used to populate the LaTeX reports
data/metadata/	Source inventory, discrepancy logs, checksums, manifest, and data dictionary
scripts/	Python-only build and analysis scripts

2 How to rebuild the package

The package is designed so that the main build command is a Python script.

```
python scripts/build_all.py --use-packaged-data
```

If you also have access to the original extracted `_Master_LCE_Docs` directory, you can regenerate the canonical datasets and the inventory directly from the source corpus:

```
python scripts/build_all.py --source-root /path/to/_Master_LCE_Docs
```

The build script regenerates:

- canonical CSV datasets;
- descriptive and discrepancy tables;
- ROC and indicator-distribution figures;
- LaTeX files and compiled PDFs;
- a manifest and SHA-256 checksums for all package files.

3 Canonical datasets

Two canonical datasets are bundled because the archive contains multiple versions of the same analyses.

- `exploratory_sample_canonical.csv` reproduces the exploratory verification sample used by the thesis and by the archived exploratory random-forest script.
- `case_study_sample_canonical.csv` reproduces the broader case-study sample whose counts and predictive metrics match the final thesis narrative.

Both canonical files are sanitised: personally identifying profile fields, links, message fragments, and contact artefacts were removed because they are not necessary for the quantitative reproduction.

4 Data dictionary

Field	Meaning
<code>listing_id</code>	Stable listing identifier inside the archived dataset
<code>source_status</code>	Archived manual status (POSITIVE, NEGATIVE, IDENTIFIED, DETECTED)
<code>ad_text</code>	Listing description text
<code>price_chf</code>	Listing price in CHF-equivalent numeric form
<code>profile_creation_year</code>	Year the seller profile was created; missing for opaque profiles
<code>content_score</code> / <code>price_score</code> / <code>profile_score</code>	Encoded operational indicators used by the original project
<code>fraud_index</code>	Sum of indicator scores used by the thesis
<code>risk_target</code>	Binary operational target: 0 = acceptable, 1 = unfavorable
<code>price_ratio</code>	Derived ratio between listing price and the model-specific market benchmark

5 Integrity and ethics notes

The source archive contains 1024 files, including legacy materials that should not be reused operationally. In particular:

- legacy scraper scripts with embedded credentials were excluded from the compendium;
- early similarity scripts with self-comparison bugs were excluded from the final analysis pipeline;
- reproduction is intentionally offline and depends on archived snapshots rather than fresh scraping.

6 Reading order

For most users, the recommended order is:

1. read the article-style reconstruction;
2. read the integrity audit to understand what was reproduced and what remains uncertain;
3. inspect the canonical datasets and the metadata tables if you need technical traceability;
4. run `python scripts/build_all.py -use-packaged-data` if you want to regenerate all outputs.